

Alicia Vanderpool

 San Francisco, CA
  john.doe@email.com
  rendercv.com
  rendercv
  rendercv

Welcome to RenderCV

RenderCV reads a CV written in a YAML file, and generates a PDF with professional typography.

Each section title is arbitrary.

You can choose any of the 9 entry types for each section.

Markdown syntax is supported everywhere. This is **bold**, *italic*, and [link](#).

Education

- | | |
|--|--|
| PhD Princeton University , Computer Science <ul style="list-style-type: none"> • Thesis: Efficient Neural Architecture Search for Resource-Constrained Deployment • Advisor: Prof. Sanjeev Arora • NSF Graduate Research Fellowship, Siebel Scholar (Class of 2022) | Princeton, NJ
Sept 2018 – May 2023 |
| BS Boğaziçi University , Computer Engineering <ul style="list-style-type: none"> • GPA: 3.97/4.00, Valedictorian • Fulbright Scholarship recipient for Graduate Studies | Istanbul, Türkiye
Sept 2014 – June 2018 |

Experience

- | | |
|--|---|
| Nexus AI , Co-Founder & CTO <ul style="list-style-type: none"> • Built foundation model infrastructure serving 2M+ monthly API requests with 99.97% uptime • Raised \$18M Series A led by Sequoia Capital, with participation from a16z and Founders Fund • Scaled engineering team from 3 to 28 across ML research, platform, and applied AI divisions • Developed proprietary inference optimization reducing latency by 73% compared to baseline | San Francisco, CA
June 2023 – present
3 years |
| NVIDIA Research , Research Intern <ul style="list-style-type: none"> • Designed sparse attention mechanism reducing transformer memory footprint by 4.2x • Co-authored paper accepted at NeurIPS 2022 (spotlight presentation, top 5% of submissions) | Santa Clara, CA
May 2022 – Aug 2022
4 months |
| Google DeepMind , Research Intern <ul style="list-style-type: none"> • Developed reinforcement learning algorithms for multi-agent coordination • Published research at top-tier venues with significant academic impact • ICML 2022 main conference paper, cited 340+ times within two years • NeurIPS 2022 workshop paper on emergent communication protocols • Invited journal extension in JMLR (2023) | London, UK
May 2021 – Aug 2021
4 months |
| Apple ML Research , Research Intern <ul style="list-style-type: none"> • Created on-device neural network compression pipeline deployed across 50M+ devices • Filed 2 patents on efficient model quantization techniques for edge inference | Cupertino, CA
May 2020 – Aug 2020
4 months |
| Microsoft Research , Research Intern <ul style="list-style-type: none"> • Implemented novel self-supervised learning framework for low-resource language modeling • Research integrated into Azure Cognitive Services, reducing training data requirements by 60% | Redmond, WA
May 2019 – Aug 2019
4 months |

Projects

FlashInfer Open-source library for high-performance LLM inference kernels - Achieved 2.8x speedup over baseline attention implementations on A100 GPUs - Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors	Jan 2023 – present
NeuralPrune Automated neural network pruning toolkit with differentiable masks - Reduced model size by 90% with less than 1% accuracy degradation on ImageNet - Featured in PyTorch ecosystem tools, 4,200+ GitHub stars	Jan 2021

Publications

Sparse Mixture-of-Experts at Scale: Efficient Routing for Trillion-Parameter Models <i>John Doe, Sarah Williams, David Park</i> 10.1234/neurips.2023.1234 (NeurIPS 2023)	July 2023
Neural Architecture Search via Differentiable Pruning <i>James Liu, John Doe</i> 10.1234/neurips.2022.5678 (NeurIPS 2022, Spotlight)	Dec 2022
Multi-Agent Reinforcement Learning with Emergent Communication <i>Maria Garcia, John Doe, Tom Anderson</i> 10.1234/icml.2022.9012 (ICML 2022)	July 2022
On-Device Model Compression via Learned Quantization <i>John Doe, Kevin Wu</i> 10.1234/iclr.2021.3456 (ICLR 2021, Best Paper Award)	May 2021

Selected Honors

- MIT Technology Review 35 Under 35 Innovators (2024)
- Forbes 30 Under 30 in Enterprise Technology (2024)
- ACM Doctoral Dissertation Award Honorable Mention (2023)
- Google PhD Fellowship in Machine Learning (2020 – 2023)
- Fulbright Scholarship for Graduate Studies (2018)

Skills

Languages: Python, C++, CUDA, Rust, Julia
ML Frameworks: PyTorch, JAX, TensorFlow, Triton, ONNX
Infrastructure: Kubernetes, Ray, distributed training, AWS, GCP
Research Areas: Neural architecture search, model compression, efficient inference, multi-agent RL

Patents

- Adaptive Quantization for Neural Network Inference on Edge Devices (US Patent 11,234,567)
- Dynamic Sparsity Patterns for Efficient Transformer Attention (US Patent 11,345,678)
- Hardware-Aware Neural Architecture Search Method (US Patent 11,456,789)

Invited Talks

- Scaling Laws for Efficient Inference — Stanford HAI Symposium (2024)
- Building AI Infrastructure for the Next Decade — TechCrunch Disrupt (2024)
- From Research to Production: Lessons in ML Systems — NeurIPS Workshop (2023)
- Efficient Deep Learning: A Practitioner’s Perspective — Google Tech Talk (2022)